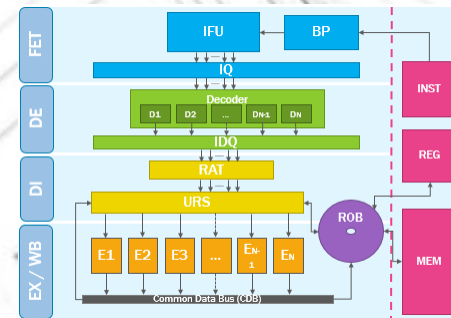


# Superscalar Processor Simulator

Felipe Galindo Sanchez

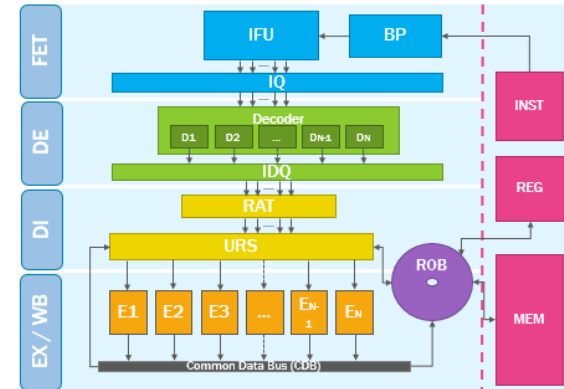
*December 2015*

*COMSM0109 – Advanced Computer*



# Superscalar Processor Simulator

- Simulator
  - Getting Started
  - Debugging
  - Compiler (Assembly to Machine code)
- Processor Architecture
  - Diagram
  - Features
  - Configurable
  - Instructions Set
- Benchmarks
  - Programs
  - Experiments



# Simulator - Getting Started



GitHub Repository

[github.com/felipegs01/processor\\_simulator](https://github.com/felipegs01/processor_simulator)

Just type **simulator.exe -help** or **./simulator --help** to list the help with all the commands and flags available, otherwise **run** is the main command.

## USAGE:

simulator run [command options] <assembly-filename>

## OPTIONS:

-s, --step-by-step	Run interactively step by step
-v, --verbose	Verbose on debug mode
-o, --output-folder	Output folder where to store debug and memory files
-c, --config-filename	Processor config filename path with the processor architecture config
--max-cycles	Maximum number of cycles to execute

Sample:

```
run samples/programs/fibonacci.asm -c samples/configs/default.config -o results/my-test --max-cycles 1000 --step-by-step -v
```

Developed in Go

<https://golang.org/>



OS Compatible



# Simulator – Debugging

## During simulation - Interactive

If the program is executed using the flag **-s** or **--step-by-step**, the state of registers and/or data memory can be seen on each cycle that is being executed.

Press the desired key and then hit [ENTER]...

- (R) to see registers memory
- (D) to see data memory
- (E) to exit and quit
- (\*) Any other key to continue

If selected **R** or **D**, the data will be displayed in the following format:

	0x00	0x04	0x08	0x0C
0x00	0x0000000A	0x0010000A	0x000C0000	0x00000000
0x10	0x000100E8	0x00000008	0x00000012	0x00000087
0x20	0x0000FF00	0x00000000	0x00D00068	0x002000A8
0x30	0x000000E8	0x0000C008	0x00000012	0x00000087
0x40	0x00000012	0x00000000	0x00100000	0x00000000

## At the end of simulation - Persisted files

At the end **six files** will be generated with details of the execution, debugging and final memory states.

The location of those output files can be selected with the flag **-o** or **--output-folder**

- **assembly.hex**: Machine code interpreted by the processor
- **memory.dat**: Final state of the data memory.
- **registers.dat**: Final state of the registers.
- **output.log**: Execution resources according to the configuration and output statistics.
- **debug.log**: Complete log for debugging purposes.
- **pipeline.dat**: Pipeline diagram of the different executed instruction stages vs execution cycles

# Simulator – Compiler (Assembly to Machine code)

This application has a builtin translator that converts *human readable assembly* instructions into *machine code*, the available instructions allowed are the ones defined on the previous [instructions](#) section.

- Only one instruction allowed per line
- Comments prefix is ;
- Comments are allowed to be on a single line or after an instruction in the same line
- It does not care about the amount of empty spaces or tabs
- Branch labels must be on a single line
- No instructions allowed to be on the same line where the branch label is declared
- Blank lines are allowed

Pre-fill data memory

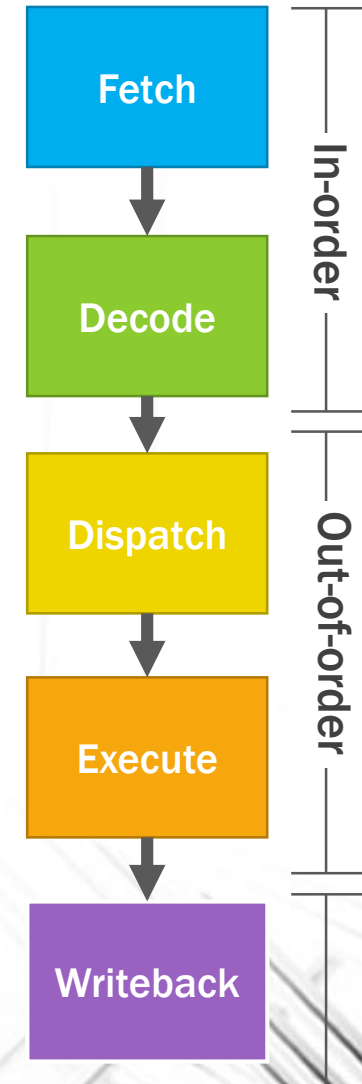
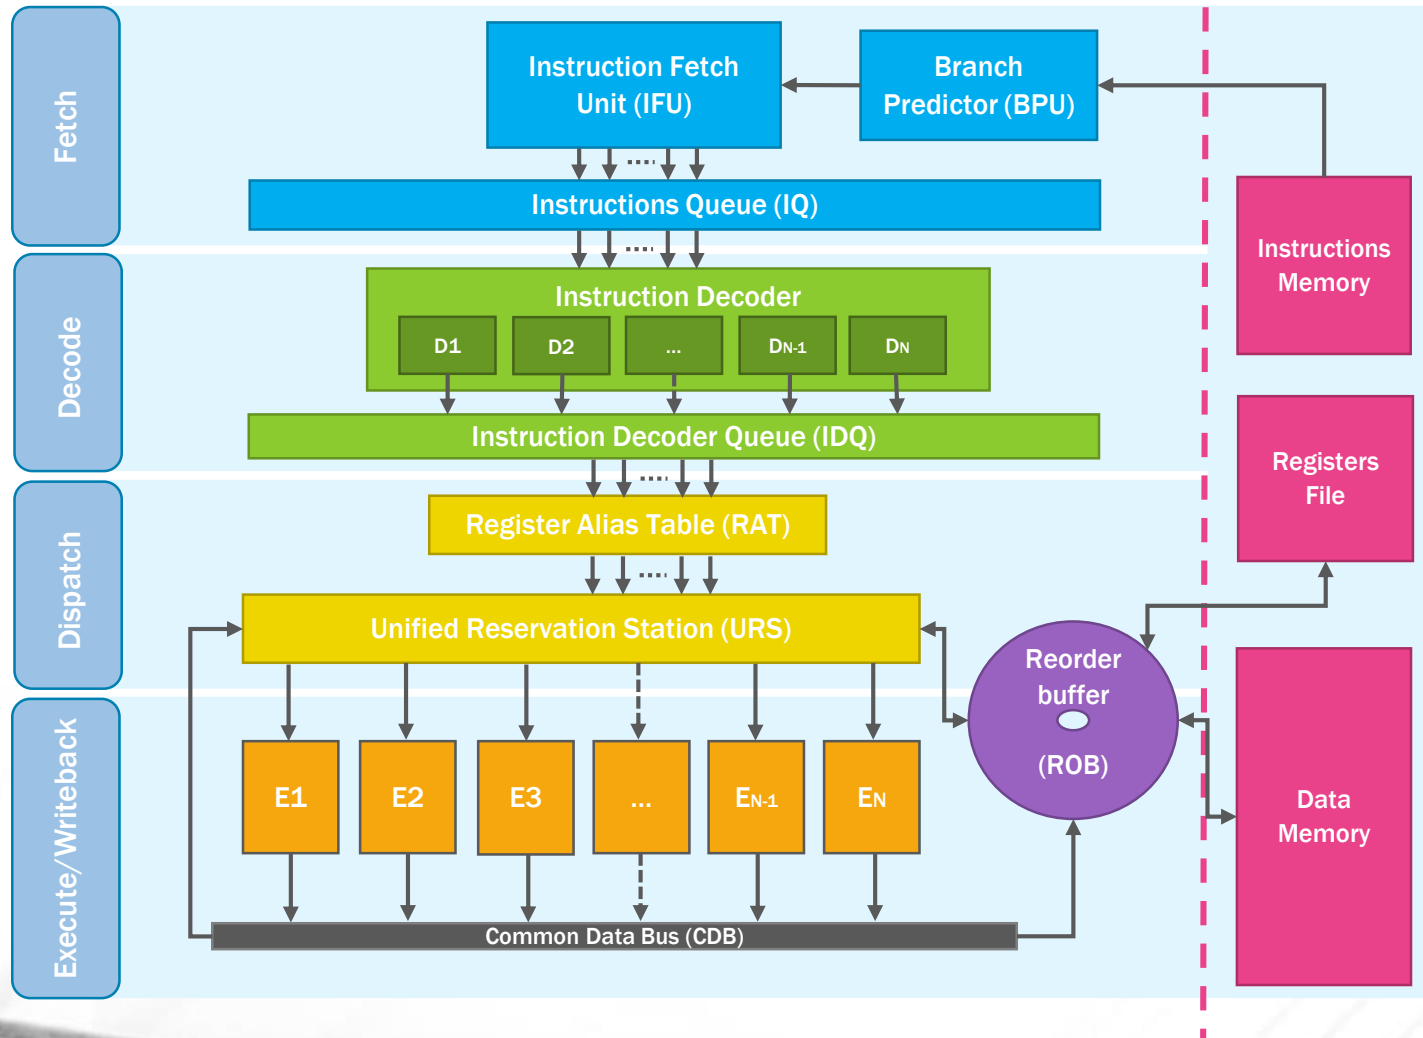
```
@0x01E0: 09 0D 23 0A 07 18 0C 06 15 0F ; Macro to pre-fill some data in memory before the program
@0x02E0: 49 2E 05 EA C6 06 3C 15 0F 41 ; is executed at an specific memory address

PROCESS_LOOP: ; Here is a label followed by an inline comment
ADDI    R1, R1, 1 ; Here is a instruction along with its operands and an inline comment
        ; Here it is an empty line which is allowed followed by an inline comment

ADD     R15, R15, R16 ; R15 += C[I]
BLT     R1, R20, PROCESS_LOOP ; Here is an instruction using a branch label followed by an inline comment
```



# Processor Architecture - Diagram



# Processor Architecture – Features (Sample)

## Overview

- 32 bits architecture
- Scalar, Pipelined or N-way superscalar
- Out-of-order execution and non-blocking issue
- 32 general purpose registers (32-bit) (used for integer & FP)
- 1 MB Instructions Memory
- 1 MB Data Memory

## Five-stage pipeline

- Fetch, Decode, Issue/Dispatch, Execute & Writeback

## Execution units (EU's)

- 2 ALU units
- 2 Load/Store units
- 1 Branch units
- 1 FPU units

## Branch Prediction

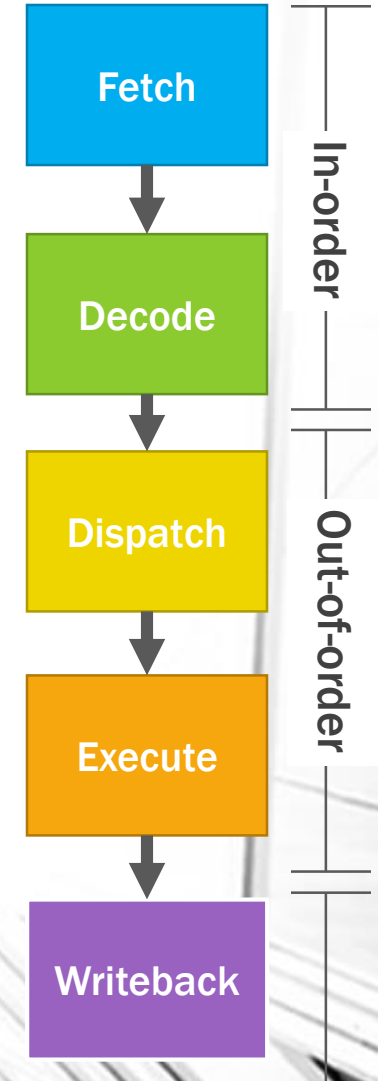
- None (Stall)
- Static: Always, Never, Forward, Backward
- Dynamic: One bit predictor, Two-bit predictor (BHT)

## Front-End Pipeline (In-order)

- Instruction Fetch Unit (IFU):
  - 16 bytes fetch on each cycle (4 instructions)
- Instruction Queue (IQ):
  - 18 instructions buffer
- Instruction Decoding Unit (IDU)
  - 4 decoding units
- Instructions Decoded Queue (IDQ):
  - 28 instructions buffer

## Execution Pipeline (Out-of-order)

- Common Data Bus (CDB)
- Register Renaming
  - Register Alias Table (RAT) with 32 entries
- Reorder buffer (ROB)
  - 32 entries
  - Up to 4 instructions written back on each cycle
- Unified reservation station (URS):
  - 128 entries
  - Up to 6 instructions dispatched on each cycle



# Processor Architecture - Configurable

## Overview

- 32 bits architecture
- **Scalar, Pipelined** or **N-way superscalar** (choose)
- Out-of-order execution and non-blocking issue
- **N** general purpose registers (32-bit) (used for integer & FP)
- **N MB** Instructions Memory
- **N MB** Data Memory

## Five-stage pipeline

- Fetch, Decode, Issue/Dispatch, Execute & Writeback

## Execution units (EU's)

- **N** ALU units
- **N** Load/Store units
- **N** Branch units
- **N** FPU units

## Branch Prediction (choose predictor)

- **None** (Stall)
- **Static:** Always, Never, Forward, Backward
- **Dynamic:** One bit predictor, Two-bit predictor (BHT)

## Front-End Pipeline (In-order)

- Instruction Fetch Unit (IFU):
  - **N** bytes fetch on each cycle (**N** instructions)
- Instruction Queue (IQ):
  - **N** instructions buffer
- Instruction Decoding Unit (IDU)
  - **N** decoding units
- Instructions Decoded Queue (IDQ):
  - **N** instructions buffer

## Execution Pipeline (Out-of-order)

- Common Data Bus (CDB)
- Register Renaming (**enabled/disabled**)
  - *Register Alias Table (RAT) with **32** entries*
- Reorder buffer (ROB)
  - **N** entries
  - *Up to **N** instructions written back on each cycle*
- Unified reservation station (URS):
  - **N** entries
  - *Up to **N** instructions dispatched on each cycle*

Configurable Parameters

### config.json

```
{  
  param1: val1,  
  param2: val2,  
  ...  
  paramN: valN  
}
```



# Processor Architecture - Instructions Set

- 32 bit instructions wide
- Instructions formats: R, I & J
- Instructions types: *Arithmetic (ALU & FPU), Load/Store, Control/Branch*

## Arithmetic/Logic (ALU & FPU)

EU	Syntax	Description	Type
ALU	add/addi Rd,Rs,Rt/C	$Rd = Rs + Rt/C$	R
	sub/subi Rd,Rs,Rt/C	$Rd = Rs - Rt/C$	R
	cmp Rd,Rs,Rt	$Rd = Rs \leq Rt$	R
	mul Rd,Rs,Rt	$Rd = Rs * Rt$	R
	shl/shli Rd,Rs,Rt/C	$Rd = Rs \ll Rt/C$	R
	shr/shri Rd,Rs,Rt/C	$Rd = Rs \ll RI/C$	R
	and/andi Rd,Rs,Rt	$Rd = Rs \& Rt/C$	R
	or/ori Rd,Rs,Rt/C	$Rd = Rs   RI/C$	R
FPU	fadd Rd,Rs,Rt	$Rd = Rs + RI$	R
	fsub Rd,Rs,Rt	$Rd = Rs - RI$	R
	fmul Rd,Rs,Rt	$Rd = Rs * RI$	R
	fdiv Rd,Rs,Rt	$Rd = Rs / RI$	R

## Formats

Type R	Format (32 bits)				
	Opcode (6)	Rd (5)	Rs (5)	Rt (5)	Not used (11)
Type I	Format (32 bits)				
	Opcode (6)	Rd (5)	Rs (5)	Immediate (16)	
Type J	Format (32 bits)				
	Opcode (6)	Address (26)			

## Load/Store

Syntax	Description	Type
lw Rd,Rs,C	$Rd = M[Rs + C]$	I
sw Rd,Rs,C	$M[Rs + C] = Rd$	I
lli Rd,C	$Rd = C$	I
sli Rd,C	$M[Rd] = C$	I
lui Rd,C	$Rd = C \ll 16$	I
sui Rd,C	$M[Rd] = C \ll 16$	I

## Control/Branch

Syntax	Description	Type
beq Rd,Rs,C	br on equal	I
bne Rd,Rs,C	br on not equal	I
blt Rd,Rs,C	br on less	I
bgt Rd,Rs,C	br on greater	I
j C	jump to C	I

# Benchmarks

REG	Registers interaction
FPO	Floating Point Operations
MEM	Memory interaction
HAZ	Data Hazards
BRN	Conditional/Unconditional branches
LOP	Loop

Program	Flavors	REG	FPO	MEM	HAZ	BRN	LOP	Special targeted features
Load/Store & ALU	Single	✓	✗	✗	✗	✗	✗	EU's throughput, Ideal N-way execution
Arithmetic Oper.	Single	✓	✓	✗	✗	✗	✗	ALU & FPU correctness
Fibonacci	Single	✓	✗	✗	✓	✓	✓	Register Renaming
Factorial	Single	✓	✗	✗	✓	✓	✗	Register Renaming
Hazards	RAW WAW    WAR	✓	✗	✗	✓	✗	✗	Reorder buffer
Loop Vectors	Forward jumps Backward jumps	✓	✗	✓	✓	✓	✓	BR Predictors
Inner Product	Single	✓	✗	✓	✓	✓	✓	Register Renaming, RS
Bubble Sort	Forward jumps Backward jumps	✓	✗	✓	✓	✓	✓	Branch Predictor, Register Renaming, RS

# Experiments

*Experiments and measurements are averaged over  $\sim 5$  executions in order to discriminate branch guesses and non-deterministic events*

# Experiment 1 – Scalar, Non-Pipelined & Superscalar

## Features targeted

Pipeline

Scalar

Superscalar

Interferences

## Hypothesis

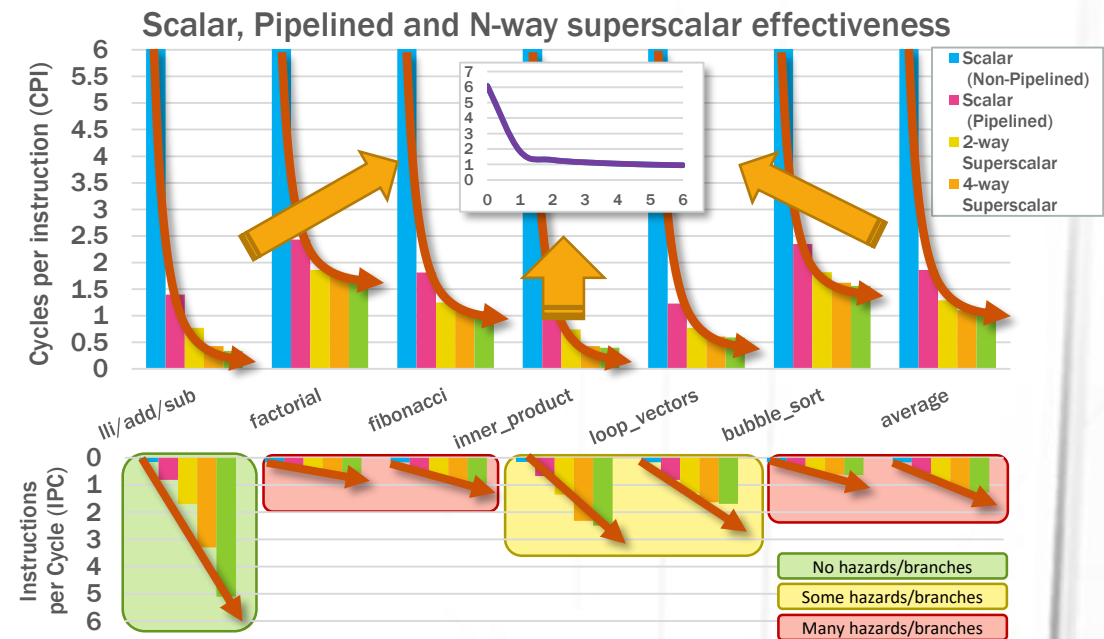
- In a **scalar non-pipelined** processor the **CPI** will be close to the **pipeline-size** (not exact due to branches)
- In a **pipelined** processor (from **1-way scalar** to **N-way superscalar**) the improvements in **CPI** will **increase**.
- In **average benchmarks/programs** including branches, hazards and non-deterministic events and the program itself it will take **one cycle per instruction**.

## Experiment

Drive all benchmarks and measure the number of cycles per instruction (CPI) for the following processor architectures:

- Scalar (non-pipelined)
- Scalar pipelined
- 2-way, 4-way, and 6-way superscalar

## Result



- **Pipelined** and **superscalar** processors definitely **improve** the **CPI/IPC**.
- **IPC** is **linearly proportional** to the **N-way** of a **superscalar**.
- The **improvement ratio** (CPI/IPC) will **depend** on **hazards** and **branches** in an executed program.
- Programs with **no hazards/branches** (e.g *lli/add/sub*) will behave close to processing **N instructions per cycle** for a **N-way superscalar** given.

# Experiment 2 – Number of Execution Units (EU's)

## Features targeted

Execution Units

## Hypothesis

- Incrementing the # of execution units will execute programs faster.
- After a number of execution units, the execution time reduced won't be that significant.

## Experiment


Drive the different available benchmarks and measure the number of cycles executed with configurations containing 3 EU's (1 branch, 1 Load/Store & 1 ALU) up to 13 EU's

Obtain the percentage of time consumed against the slowest configuration (3 EU's)

$$\% \text{ exec time (\#EU's)} = \frac{\# \text{ cycles executed (\#EU's)}}{\# \text{ cycles executed (3 EU's)}}$$

Find the # EU's where the % savings is not significant anymore.

## Result



- From 3 up to 6 EU's, programs are **11% faster per each EU**
- From 7 up to 13 EU's, programs are **1.57% faster per each EU**
- Therefore it's concluded that incrementing the # of EU's will improve the execution time with a **significant % of savings up to ~6 EU's**, but more than that, the % of savings won't be that significant

# Experiment 3 – Branch Predictors and miss-prediction

## Features targeted

Fetch Unit

Branch Predictors

## Hypothesis


- Branch predictors improves the execution time of programs
- Static predictors can perform good or bad depending on how programs are coded with a high variability
- Dynamic predictors perform better no matter which way programs are coded with a lower variability.

## Experiment

Drive benchmark programs coded with two flavors: *backward jumps* and *forward jumps*, in order to benchmark the different results with different programming ways

Compare benchmarks and flavors with the different branch predictors from *none*, *static* to *dynamic* and measure the percentage of miss prediction -> (the more miss prediction, the more time it will take to execute a program)

## Result



- We can see that **static predictors** have a bigger miss prediction variability range depending on how the program is coded (conditional branches)  
Approx. **50% miss prediction** with **+/- 50% variability**
- **Dynamic predictor** performed better no matter what flavor was executed  
Approx. **20% miss prediction** with **+/- 10% variability**



# Experiment 4 – Reorder Buffer Size

## Features targeted

Reorder Buffer

Reservation Station

## Hypothesis

- Due to ROB's capacity, some instructions will be stalled on dispatching if the ROB is full (meaning there are pending results to be committed), therefore incrementing the # of entries in the buffer, the speed will be incremented.
- The speed increment will stop at a certain # of ROB entries.

## Experiment


Drive the different available benchmarks and measure the number of cycles executed with configurations containing a ROB starting from 1 entry up to 60 entries.

Obtain the speed as a function of the # of entries against the slowest (1-entry ROB)

$$\text{speed (ROB size)} = \frac{\# \text{ cycles executed (1-entry ROB)}}{\# \text{ cycles executed (ROB size)}} \times$$

Find the # of ROB entries where the increment stops.

## Result



- From a **1-entry** ROB to a **~25-entry** ROB the speed will be incremented approximately **1.24X** or **24%**. (**1.1% or 1.1X** for each added **register**)
- ROB's with more than **~25 entries**, (unless programs with mixed hazards provoking to have more than ~25 dispatched instructions and pending to be committed) **won't make any improvement** at all.

# Experiment 5 – Register re-naming using RAT

## Features targeted

Register Re-naming

Register Alias Table (RAT)

## Hypothesis


- Using *register re-naming* will avoid stall WAR and WAW register dependencies
- Once *register re-naming* enabled, the RAT size will improve the execution speed until an approximated RAT size when the speed improvement won't be that significant.

## Experiment

- 1) Drive RAW, WAR and WAW hazards with and without *register re-naming* enabled. Then with the pipeline diagram determine if the dependency was or was not halted.
- 2) Drive available benchmarks and measure percentage of execution time against *register re-naming* disabled

$$\% \text{ exec time } (RAT_{size}) = \frac{\# \text{ cycles executed } (RAT_{size})}{\# \text{ cycles executed } (RR_{disabled})}$$

## Result



- Using **register re-naming** definitely improves the performance, **avoiding halting** on **WAR** and **WAW** dependencies by using other physically registers in the **Register Alias Table (RAT)**.
- Using a **3 to 7-entries RAT**, programs are faster approximately in **40%**. (**5.7%** faster for **each** added **entry** to the RAT)
- Using a **7 to 30-entries**, programs can be **18%** faster (not that significant). **0.8%** faster for **each** added **entry** to the RAT)